

MY SERVER IS LYING ABOUT MY DISKS

Here is an interesting problem I can up against again recently. On a linux server, when you want to see your disk usage, you can do this..

```
# df -h
Filesystem      Size  Used Avail Use% Mounted on
/dev/sda1        99M   12M   82M   13% /boot
tmpfs            1014M    0 1014M    0% /dev/shm
/dev/sdb1       493G  319G  149G   69% /remotelog
```

..and we see that "/remotelog" is 69% used. End of Story. But, what if you decided you wanted to just check where the usage occurs, so you do..

```
# du -csh /remotelog/
155G    /remotelog/
155G    total
```

Suddenly you get two very different answers. Engage HCM (Headless Chicken Mode), run around the room screaming either "My server is haxored" or "I have gone mad" or a combination of the two. Once you got that out your system disengage HCM and lets take a look at this. You see, neither of the previous two explanations are right. The answer is actually fairly simple. *df* and *du* both report on disk usage, but they both (and I am simplifying here) get there data from different places. *du* looks at what is on the disk (like adding up the file sizes in a *ls*) while *df* looks at what the system is doing. While they should be the same, they can differ. On this system for example, lets do a *ls* (a glorious tool)..

```
# ls -l | grep remotelog
-rw-r--r-- 1 root root 12106 2010/09/30 10:17 remotelog/2010/09/30/10.172.80.17.log (deleted)
-rw-r--r-- 1 root root 15386 2010/09/30 10:17 remotelog/2010/09/30/choprly1.log (deleted)
-rw-r--r-- 1 root root 94709719760 2010/10/01 10:17 remotelog/2010/10/01/10.172.3.172.log
-rw-r--r-- 1 root root 1568 2010/10/01 10:17 remotelog/2010/10/01/10.172.2.24.log
-rw-r--r-- 1 root root 1918 2010/10/01 10:17 remotelog/2010/10/01/192.168.50.127.log
-rw-r--r-- 1 root root 1922 2010/10/01 10:17 remotelog/2010/10/01/10.172.103.27.log
-rw-r--r-- 1 root root 1680 2010/10/01 10:17 remotelog/2010/10/01/10.172.2.25.log
-rw-r--r-- 1 root root 5552 2010/10/01 10:17 remotelog/2010/10/01/10.172.80.17.log
-rw-r--r-- 1 root root 11603 2010/10/01 10:17 remotelog/2010/10/01/choprly1.log
-rw-r--r-- 1 root root 176674191060 2010/10/01 10:17 remotelog/2010/09/30/10.172.3.172.log (deleted)
```

And the penny (or locale-specific minimum denomination currency unit) drops. The 'system' is still busy with some files on that disk. The files may be gone, but for whatever reason the files are still in use by the system. This is why *du* and *df* are giving different answers. So to fix this, a simple..

```
/etc/init.d/rsyslogd restart
```

..and wait a bit for the old process to wrap up. Then when we do a *df* we get..

```
# df -h
Filesystem      Size  Used Avail Use% Mounted on
/dev/sda1        99M   12M   82M  13% /boot
tmpfs           1014M    0 1014M   0% /dev/shm
/dev/sdb1       493G  156G  312G  34% /remotelog
```

..and there we go, our server utilities are back to agreeing with one another.

Final Words

This little hiccup is kinda interesting in it's own way, but to me what it highlights more is that (1) faulty understanding can lead to faulty conclusions and (2) it is important to know not only how the tools work but why they work. Go play, have fun and learn.